

Deep Learning

Building and improving
deep learning architectures




Dr Chris Willcocks

Department of Computer Science



Aims

1. How to design the mind?
2. State-of-the-art components



Lecture Overview

Today's lecture

- Convolutional layers
 - Convolution arithmetic
 - Stride, padding, and dilations
 - Common patterns
 - Transpose convolutions
- Pooling layers
 - Mean pool, max pool, adaptive pooling, interpolation
- Regularisation techniques
 - L1, L2
 - Dropout
 - Normalisation layers
- Residual learning

Materials:



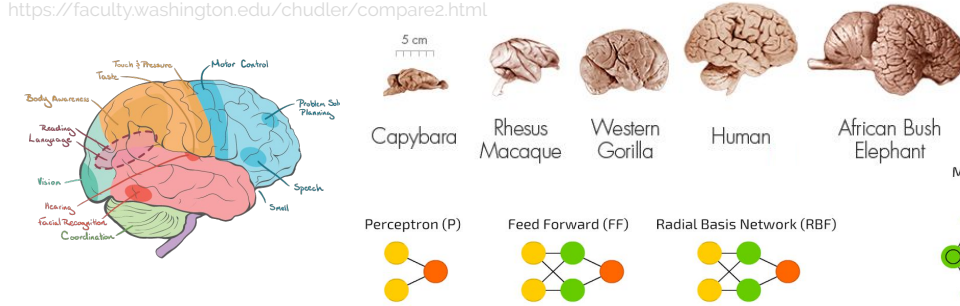
<https://github.com/cwkk/ml-materials>



<https://colab.research.google.com/gist/cwkk/1a4da17e2e9c1081f93631f5b51e5bae/convnet.ipynb>

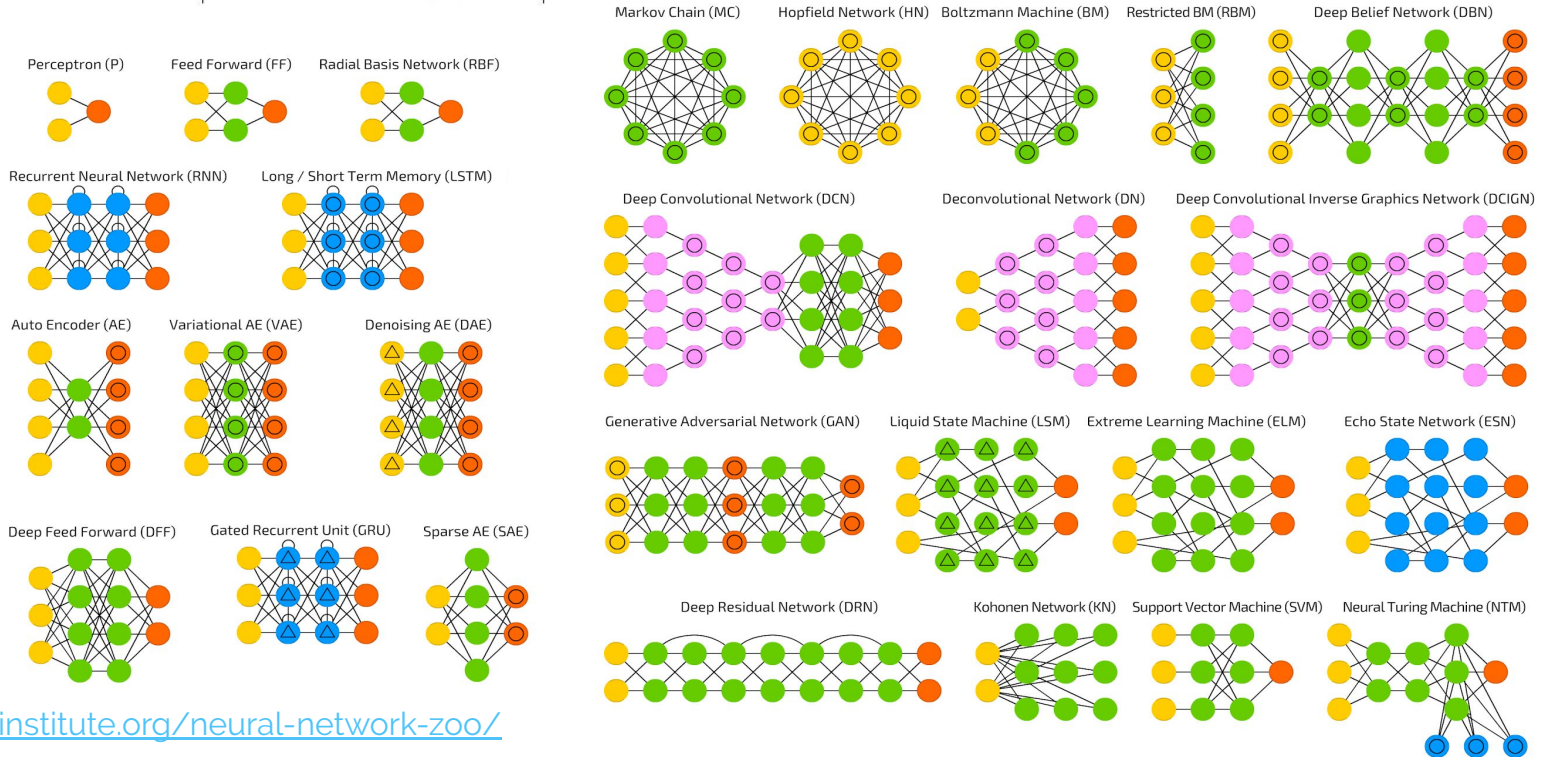
Building and improving architectures

<https://faculty.washington.edu/chudler/compare2.html>




- Architecture is very important to function
 - We'll look today at main components

- Backfed Input Cell
- Input Cell
- Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- Different Memory Cell
- Kernel
- Convolution or Pool



Fully-connected layer properties

- Big inputs =
lots of parameters
- Expensive
 - Memory
 - Processing
- Fixed



```
(base) chris@chris-office ~ master • ipython
Python 3.7.1 (default, Oct 23 2018, 19:19:42)
Type 'copyright', 'credits' or 'license' for more information
IPython 7.1.1 -- An enhanced Interactive Python. Type '?' for help.

In [1]: import torch

In [2]: import torch.nn as nn

In [3]: x = torch.zeros(1, 3, 1024, 1024)


In [4]: l1 = nn.Linear(1024*1024*3, 512)

In [5]: l1.weight.view(-1).size()
Out[5]: torch.Size([1610612736])
```

- 6 GB of memory if float 32-bit
- Just 1 layer
- And that's only 512 output features...

Convolutional layers (1D, 2D, 3D)

Kernel slides a window over input...



learn these weights



$$* \begin{bmatrix} -1, 0, 1 \\ -2, 0, 2 \\ -1, 0, 1 \end{bmatrix} =$$



- <http://setosa.io/ev/image-kernels/> - interactive 2D demo
- https://github.com/vdumoulin/conv_arithmetic - detailed diagrams
- <https://arxiv.org/pdf/1603.07285.pdf> - great paper (diagrams above)

Convolution arithmetic




- More advanced usage helps us control the shape of the architecture
 - Kernel size, stride, padding, and dilations

```
nn.Conv2d(in_channels, out_channels, kernel_size, stride=2, padding=1, dilation=2, groups=1, bias=True)
```

e.g. on
first layer

{ R
G
B

Ridge 90°
Ridge 33°
Ridge 45°
Blob
Texture
...



Common Patterns

- Input: (N, C_{in}, L_{in}) ← Spatial dimension(s)
- Output: (N, C_{out}, L_{out}) where

$$L_{out} = \left\lfloor \frac{L_{in} + 2 \times \text{padding} - \text{dilation} \times (\text{kernel_size} - 1) - 1}{\text{stride}} + 1 \right\rfloor$$

```
nn.Conv2d(in_channels=3, out_channels=14, kernel_size=3, stride=1, padding=1, dilation=1)
```

$[B \times 3 \times 64 \times 64] \rightarrow [B \times 14 \times 64 \times 64]$ - don't change spatial dimension

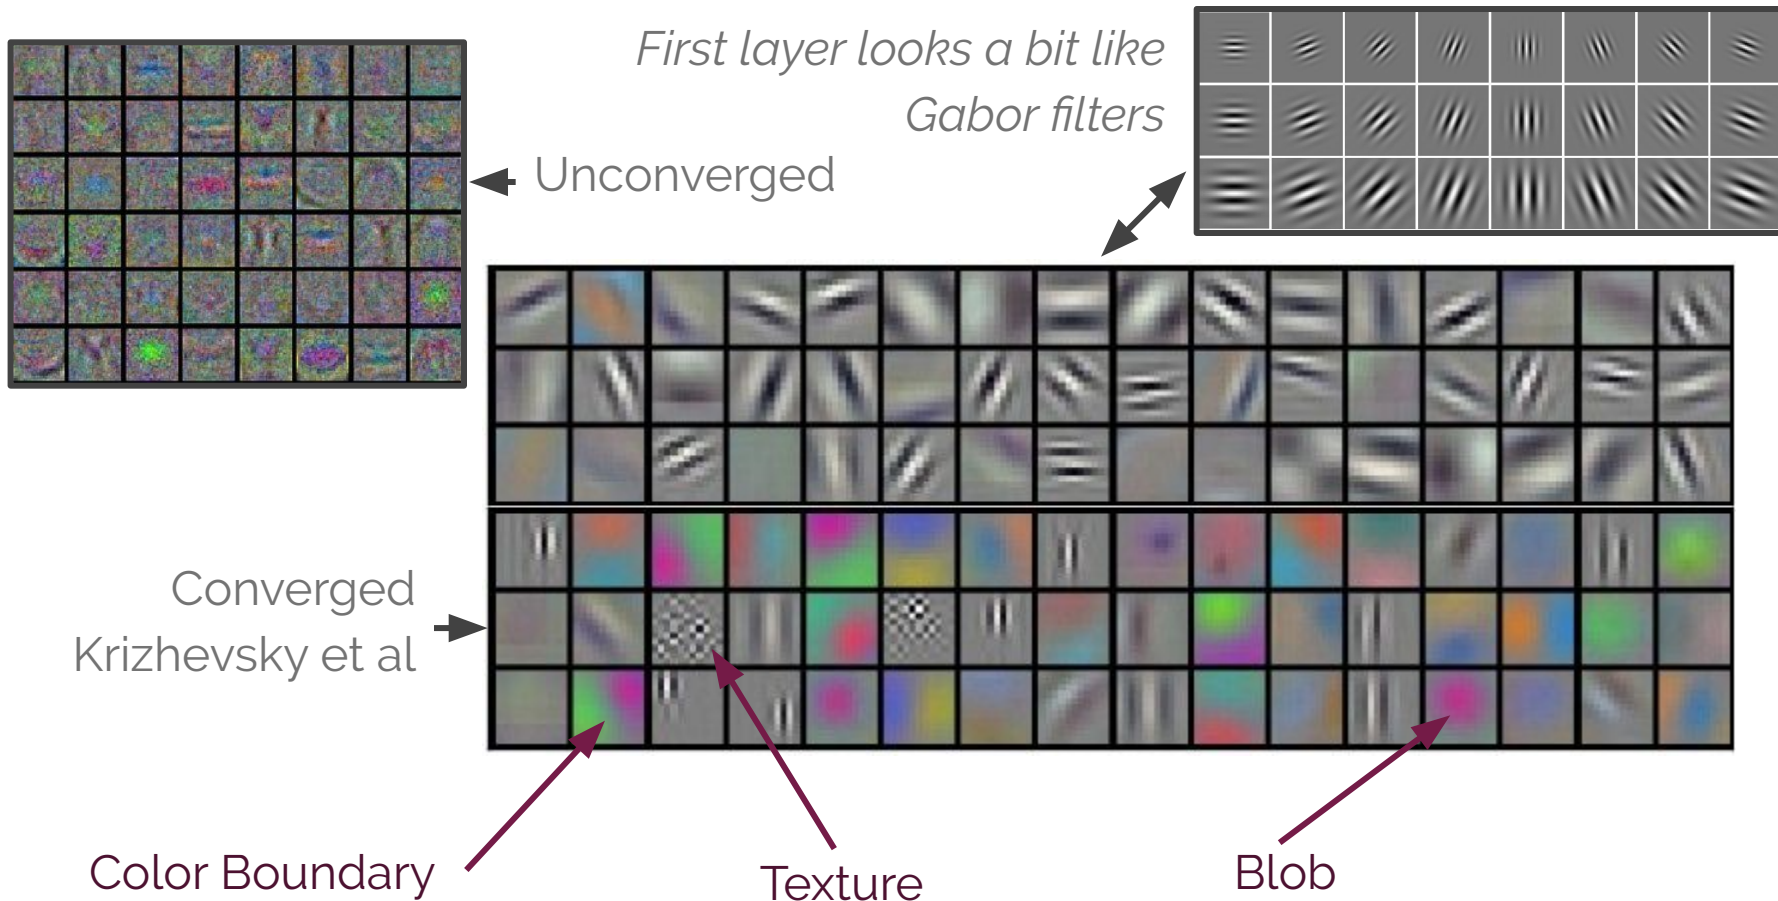
```
nn.Conv2d(in_channels=3, out_channels=1, kernel_size=1, stride=1, padding=0, dilation=1)
```

$[B \times 3 \times 64 \times 64] \rightarrow [B \times 1 \times 64 \times 64]$ - don't change spatial dimension

```
nn.Conv2d(in_channels=3, out_channels=14, kernel_size=4, stride=2, padding=1, dilation=1)
```

$[B \times 3 \times 64 \times 64] \rightarrow [B \times 14 \times 32 \times 32]$ - half spatial dimension

First layer weights respond to edges, color variation, texture...




Transpose convolutions

- Input: (N, C_{in}, L_{in})
- Output: (N, C_{out}, L_{out}) where

$$L_{out} = (L_{in} - 1) \times \text{stride} - 2 \times \text{padding} + \text{kernel_size} + \text{output_padding}$$

- “The gradient of a convolution with respect to its input”
- Useful for going back up (increasing spatial dimension)
 - E.g. generators and decoder components



```
nn.Conv2d(in_channels=3, out_channels=14, kernel_size=4, stride=2, padding=1, dilation=1)
```

$[B \times 3 \times 64 \times 64] \rightarrow [B \times 14 \times 128 \times 128]$ - *doubled the spatial dimension*

Other ways to change spatial dimension

- In contrast to fully-convolutional (recommended)
- Other ways to change the spatial dimension
 - Mean pool
 - Max pool
 - Interpolation (can go up or down)
- Adaptive pooling (can go up or down)
 - The kernel can be altered based on the size of the input
- Common approaches
 - Can implemented in the forward pass using **nn.functional** interface as these methods don't require parameters

```
nn.AdaptiveAvgPool2d((5,71))
```


$[B \times 3 \times 64 \times 64] \rightarrow [B \times 3 \times 5 \times 71]$ -
output is fixed

```
nn.AdaptiveAvgPool2d((1,1))
```


$[B \times 3 \times 27 \times 19] \rightarrow [B \times 3 \times 1 \times 1]$ -
good for last layer

Regularisation


- “Regularization is any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.”
 - Goodfellow’s book, Section 5.2.2
- Approaches
 - Add “error”, e.g.:
 - Dataset augmentation
 - Dropout
 - Constrain manifold, e.g.:
 - Smooth the optimization landscape



Underfitting




Good fit



Overfitting



Augmentation

- A form of adding prior knowledge to the model, e.g. fundamentally these are all recognizable images of dogs →
- Examples:
 - Random rotations
 - Random horizontal flips
 - Random blur
 - Random noise
- Would 180 degree rotations be suitable data augmentation for MNIST?



L1 and L2 regularisation

Recall Hebbian
theory and synaptic
plasticity...



It'd be nice if some of the weights could $\rightarrow \mathbf{o}$

L1 regularizer is not differentiable everywhere but allows sparsity

$$\begin{aligned}\mathcal{L}' &= \mathcal{L} + \lambda \sum_i |w_i| \\ &= \mathcal{L} + \lambda \|\mathbf{w}\|_1\end{aligned}$$


L2 regularizer is convex and differentiable everywhere

$$\begin{aligned}\mathcal{L}' &= \mathcal{L} + \lambda \sum_i w_i^2 \\ &= \mathcal{L} + \lambda \|\mathbf{w}\|_2^2\end{aligned}$$


Stochastic Regularisation

Dropout & Label Noise


- Each hidden unit is set to zero with some probability (e.g. 0.2)
- Cannot rely on any one weight.
 - Spreads out its weights
 - Shown formally to have similar effect to l2 regularisation



(a) Standard Neural Net



(b) After applying dropout.



Normalisation


Batch Normalization: Accelerating Deep Network Training by ...

<https://arxiv.org> > cs ▼

by S Ioffe - 2015 - Cited by 8392 - Related articles

11 Feb 2015 - Applied to a state-of-the-art image classification model, **Batch Normalization** achieves the same accuracy with 14 times fewer training steps, ...

- Batch normalisation in particular is a very important technique
- Recently it seems main results come from smoothing optimisation landscape rather than preventing covariance shift
 - <https://arxiv.org/abs/1805.11604>
- Other normalisation layers (useful in different scenarios):



Different normalisation layers (figure from Group Norm paper: <https://arxiv.org/pdf/1803.08494.pdf>). N=batch axis, C=channel axis, H,W = spatial axes. Blue voxels are normalized to have same mean and variance, computed by aggregating values of those voxels.



Deep Residual Learning

- What if we want to go really deep?
 - Vanishing gradient problem
- Residual blocks

Deep Residual Learning


<https://arxiv.org> > cs ▼

by K He - 2015 - Cited by 17392 -
10 Dec 2015 - We present a **residual**
evidence showing that these **residual**



(a) original

Example Convnet Implementation



```
class ResidualBlock(nn.Module):
    def __init__(self, in_features):
        super(ResidualBlock, self).__init__()

        conv_block = [ nn.Conv2d(in_features, in_features, 3, stride=1, padding=1, bias=False),
                        nn.BatchNorm2d(in_features),
                        nn.ReLU(inplace=True),
                        nn.Conv2d(in_features, in_features, 3, stride=1, padding=1, bias=False),
                        nn.BatchNorm2d(in_features) ]

        self.conv_block = nn.Sequential(*conv_block)

    def forward(self, x):
        return torch.relu(x + self.conv_block(x))
```

Take away points

- Convolutions are great when you have **spatial** or **temporal coherence**
- Enforcing **smoothness** (where applicable) generally improves stability and test accuracy
 - Smoothness as regularizer
 - Smoothness during data transformation
- Regularize if you are overfitting
- The shape of the architecture is important to the application
- Residual layers help go deep without vanishing gradients

