

Machine Learning

Reinforcement Learning

Dr Chris Willcocks
Department of Computer Science



Lecture Overview

Recap

- GANs, AEs, Latent spaces, UNet...
 - Mostly differentiable problems up until now
- Synaptic plasticity

Today's lecture: Reinforcement Learning

“Reinforcement”

- “The action or process of reinforcing or strengthening”
- “The process of encouraging or establishing a belief or pattern of behaviour”

Further Reading: [Reinforcement Learning: An Introduction](#) [Lilian Weng](#) [David Silver](#)
(if you're interested in learning more about this field)





Illustration by Joshua Seong, Verywell

Reinforcement Learning

Given an agent in an unknown environment, learn how to take **smart actions** such as to **maximize cumulative rewards**




- There is **no supervisory signal**, only a notion of reward R_t at step t
- Feedback is **not immediate**
- **Time** is important (sequential, data is not *iid*)
- Agents actions affect subsequent environment state data



Examples

Control behaviour of complex systems

- Defeat champion at Go
- Do stunt manoeuvres in drones
- Beat humans at video games
- Find most efficient robot walking strategy over complex terrain
- Control power stations
- Self-driving cars
- Robotic cook dinner
- Investing in stock market
-



Learning to drive and flip pancakes



Wayve: Learning to drive in a day



Learning to flip pancakes

Key Concepts


Interaction between **agent** and **environment** through sequential time steps $t = 1, 2, \dots, T$. The agent *learns* about the environment, it learns the optimal **policy** and makes decisions about how to choose the next action.

Interaction sequence is described by an **episode (trajectory)** consisting of **states**, **actions**, and **rewards** at given time steps: S_t, A_t and R_t :

$$S_1, A_1, R_2, S_2, A_2, \dots, S_T$$

There are many categories of RL algorithm, generally they are:

- **Model-based:** The agent learns a model of the environment
- **Model-free:** The agent learns which actions to take without an environment model
- **On-policy:** Learning on the job
- **Off-policy:** Learning optimal policy independent of current policy being executed



Key Concepts

Modelling the environment's transition and reward

The **model** describes the environment

1. Transition probability function P
2. Reward function R

Transition steps: $\langle s, a, s', r \rangle$



<https://www.placesyoullsee.com/25-most-dangerous-hiking-trails-in-the-world/>

$$P(s', r | s, a) = \mathbb{P}[S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a]$$

$$R(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} P(s', r | s, a)$$

Key Concepts

Policy and value

- The agent's “policy” π describes which action it should take in a given state.

It can be deterministic: $a = \pi(s)$

or stochastic:

$$a = \pi(a|s) = \mathbb{P}[A_t = a|S_t = s]$$



- The value function estimates the **future reward** for a state or an action. The future reward is called the **return**:

$$G_t = R_{t+1} + \gamma R_{t+2} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Key Concepts

Value functions (continued)

- The value function estimates the **future reward** for a **state** or an **action**.

The **state**-value for a state s is the expected return at time t

$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

The **action**-value (Q-value, or “quality”) for a state-action pair is similarly:

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

The **advantage** is the difference between the action-value and state-value:

$$A_{\pi}(s, a) = Q_{\pi}(s, a) - V_{\pi}(s)$$

Key Concepts

Optimal value and policy

The **optimal value** function $V_*(s)$ produces the maximum return

$$V_*(s) = \max_{\pi} V_{\pi}(s), Q_*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

The **optimal policy** π_* is the policy that corresponds to the optimal value function

$$\pi_* = \arg \max_{\pi} V_{\pi}(s), \pi_* = \arg \max_{\pi} Q_{\pi}(s, a)$$

where $V_{\pi_*}(s) = V_*(s)$ and $Q_{\pi_*}(s, a) = Q_*(s, a)$



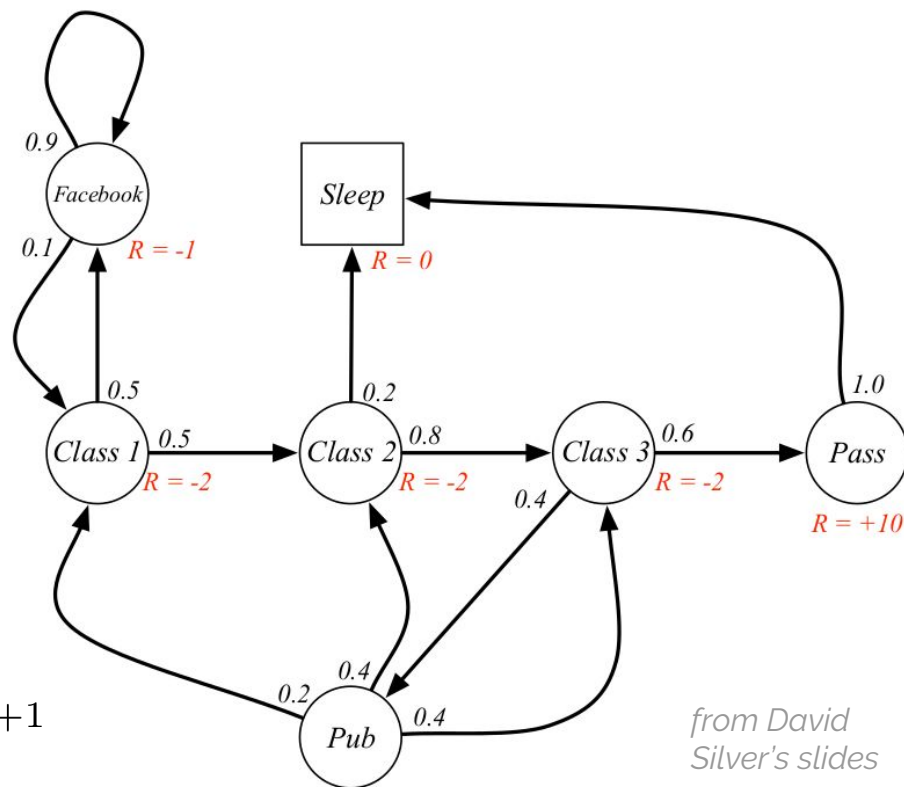
Markov Decision Process

Most *reinforcement* models as **Markov decision processes** $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, P, R, \gamma \rangle$

1. Set of states \mathcal{S}
2. Set of actions \mathcal{A}
3. Transition function $P(s'|s, a)$
4. Reward function $R(s, a, s')$
5. Discount factor $\gamma \in [0, 1]$

- Initial state s_0
- Discount factor $\gamma \in [0, 1]$ for return G_t recall

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



from David
Silver's slides


Markov States

"The future is independent of the past given the present"

A state \mathcal{S}_t is *Markov* if and only if

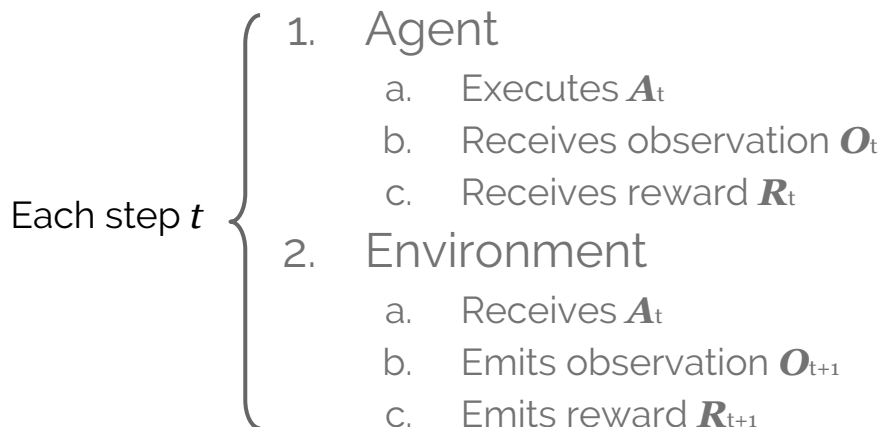
$$\mathbb{P}[\mathcal{S}_{t+1} | \mathcal{S}_t] = \mathbb{P}[\mathcal{S}_{t+1} | \mathcal{S}_1, \mathcal{S}_2, \dots, \mathcal{S}_t]$$

- The state captures all information from the history
- Once the state is known, the history can be thrown away
- The state is a sufficient statistic of the future




How to make decisions?

- Select actions to *maximise total future reward*
- Actions may have long term consequences
- Reward may be delayed
- Sometimes sacrifice of immediate reward is necessary for long-term reward
 - Refuelling a racing car

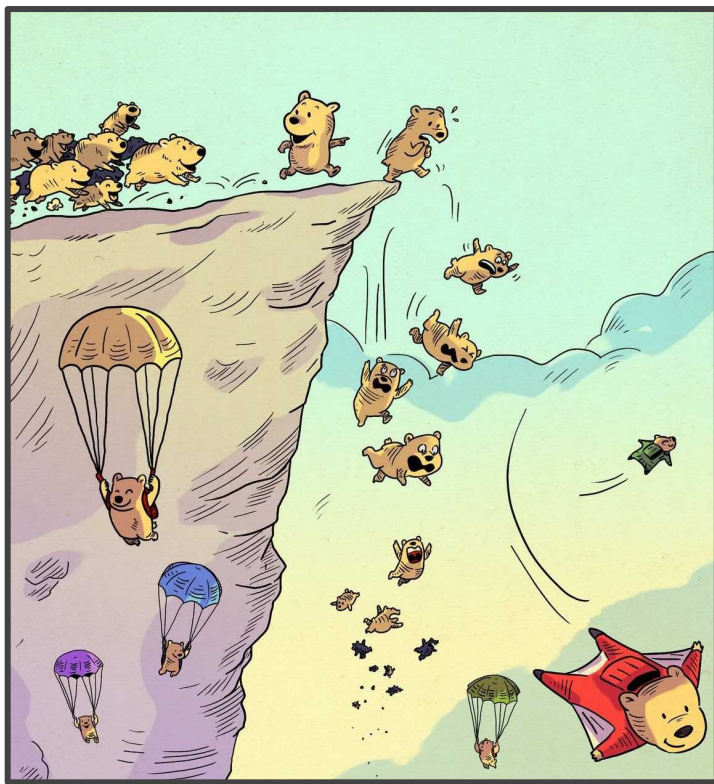


Partial observability

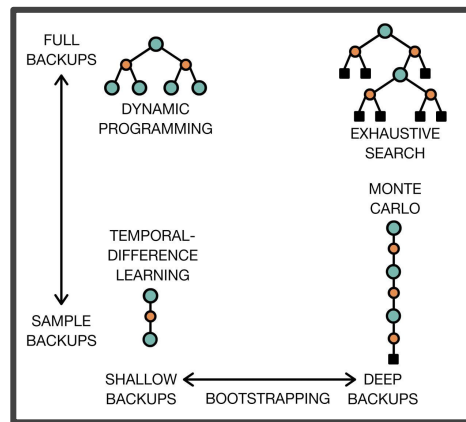


from slides by David Silver

Common Approaches



- Dynamic Programming
- Monte-Carlo methods
- Temporal-Difference Learning
- Policy Gradients
- Evolution Strategies



Common Approaches


Dynamic Programming

If we have complete information about the environment, and know $P(s'|s, a)$ and $R(s, a, s')$, we can use [Dynamic Programming](#) to directly solve MDP's by applying the [Bellman Optimality Equations](#), which show that:

$$\begin{aligned} V(s) &= \mathbb{E}[G_t | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s] \end{aligned}$$

We can then iteratively evaluate the value function and improve the policy

$$\pi_0 \xrightarrow{\text{eval}} V_{\pi_0} \xrightarrow{\text{refine}} \pi_1 \xrightarrow{\text{eval}} V_{\pi_1} \xrightarrow{\text{refine}} \pi_2 \xrightarrow{\text{eval}} \dots \xrightarrow{\text{refine}} \pi_* \xrightarrow{\text{eval}} V_*$$



Common Approaches

Monte-Carlo Methods

Monte-Carlo (MC) methods simply learn through experience without modelling the environment dynamics.

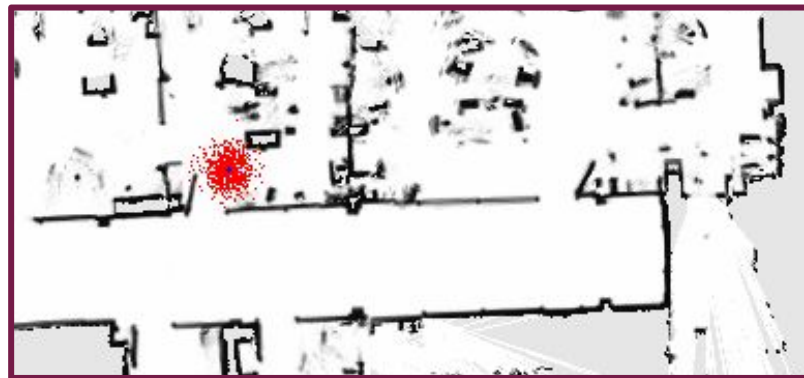
Typical approach:

1. Improve the policy greedily with respect to current value function

$$\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$$

2. Generate new trajectory with new policy π with ϵ -greedy policy.
3. Estimate the new state-action return Q , where

$$q_{\pi}(s, a) = \frac{\sum_{t=1}^T (1[S_t = s, A_t = a] \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1})}{\sum_{t=1}^T 1[S_t = s, A_t = a]}$$



Common Approaches

Temporal-Difference Learning

Temporal-Difference (TD) learning is where we slowly move the value function $V(S_t)$ towards $R_{t+1} + \gamma V(S_{t+1})$ which is the TD target by some learning rate hyper parameter α

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

TD Target

- No reliance on rewards
- No complete returns (incomplete episodes/trajectories)

Bootstrapping update targets with regard to existing estimates, rather than relying on actual rewards and complete returns.

Wed



16° 3°

Thu



11° 4°

Fri



12° 6°

Sat



12° 5°

Sun



11° 4°



Common Approaches

Temporal-Difference Learning: Q-Learning

The previous slide also applies for estimating the return of state-action pairs:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

But what if we did an **off-policy** approach, instead of sampling SARSA with ϵ -greedy, estimating Q_* out of the best values independent of the current policy?



$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

Memorizing $Q_*(.)$ for all state-action pairs is large/expensive. So we use a function approximator $Q(s, a; \theta)$ aka a **Deep Neural Network**.

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right]$$

Common Approaches

TD Learning: Deep Q-Learning

[PDF] Playing Atari with Deep Reinforcement Learning - University of ...

<https://www.cs.toronto.edu/~vmnih/docs/dqn.pdf> ▼

by V Mnih - Cited by 2031 - Related articles

However **reinforcement learning** presents several challenges from a **deep** ...

This **paper** demonstrates that a convolutional neural network can overcome

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for




Achievement unlocked! You can now play Atari games!

Common Approaches

Policy Gradients

Previously, we've tried to learn the state/action value and choose actions.

What if we want to learn the policy function $\pi(a|s; \theta)$ with respect to some parameters θ ?



How would we train this?

Common Approaches

Policy Gradients

First, we define the reward function

$$\mathcal{J}(\theta) = V_{\pi_{\theta}}(S_1) = \mathbb{E}_{\pi_{\theta}}[V_1]$$

To do gradient ascent, we need the partial derivative of the reward function with respect to the parameters, e.g. *numerically*

$$\frac{\partial \mathcal{J}(\theta)}{\partial \theta_k} \approx \frac{\mathcal{J}(\theta + \epsilon u_k) - \mathcal{J}(\theta)}{\epsilon}$$

where

$$\nabla \mathcal{J}(\theta) = \mathbb{E}_{\pi_{\theta}}[\nabla \ln \pi(a|s, \theta) Q_{\pi}(s, a)]$$

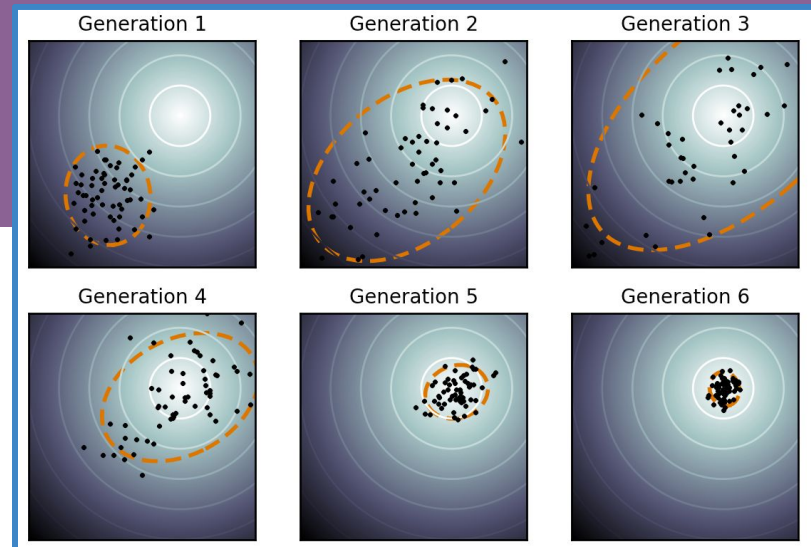
see **REINFORCE** and **Actor-Critic**, which are full learning algorithms that do this.
[Link to an excellent write-up by Lilian Weng](#) on policy-gradient algorithms.

Common Approaches

Evolution Strategies

If we satisfy:

1. Solutions able to freely interact with the environment and see what they do
2. Fitness for any solution can be evaluated



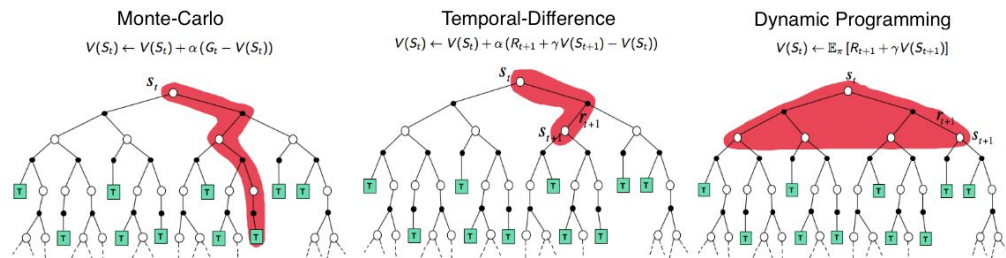
... then we can use **evolution strategies** (e.g. [CMA-ES](#), Genetic Algorithm) using a non MDP-based approach without value approximation.

We assume a prior distribution over the policy parameters θ (e.g. a multivariate Gaussian) and sample the gradient accordingly

$$\nabla_{\theta} \mathbb{E}_{\epsilon \sim N(0, I)} F(\theta + \sigma \epsilon) = \frac{1}{\sigma} \mathbb{E}_{\epsilon \sim N(0, I)} [F(\theta + \sigma \epsilon) \epsilon]$$

Take away points

- Reinforcement learning is a huge field still in its infancy...
 - There are lots of ways to build RL agents
 - There are lots of constraints in different applications
 - Quite expensive (compute, training examples)
 - **Exploitation vs exploration**
- We've only dipped our feet into it!



...where we're currently at

1. Continuous form of policy reward function

$$\begin{aligned}\mathcal{J}(\theta) &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi(a|s; \theta) Q_{\pi}(s, a) \\ \nabla \mathcal{J}(\theta) &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \nabla \pi(a|s; \theta) Q_{\pi}(s, a) \\ &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi(a|s; \theta) \frac{\nabla \pi(a|s; \theta)}{\pi(a|s; \theta)} Q_{\pi}(s, a) \\ &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi(a|s; \theta) \nabla \ln \pi(a|s; \theta) Q_{\pi}(s, a) \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla \ln \pi(a|s; \theta) Q_{\pi}(s, a)]\end{aligned}$$

2. Lerp value-return for TD-learning

$$\begin{aligned}V(S_t) &\leftarrow (1 - \alpha)V(S_t) + \alpha G_t \\ V(S_t) &\leftarrow V(S_t) + \alpha(G_t - V(S_t)) \\ V(S_t) &\leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))\end{aligned}$$

Appendix

Notation 1/3

s	state
a	action
\mathcal{S}	set of all nonterminal states
\mathcal{S}^+	set of all states, including the terminal state
$\mathcal{A}(s)$	set of actions possible in state s
\mathcal{R}	set of possible rewards
t	discrete time step
T	final time step of an episode
S_t	state at t
A_t	action at t
R_t	reward at t , dependent, like S_t , on A_{t-1} and S_{t-1}
G_t	return (cumulative discounted reward) following t

Appendix

Notation 2/3



π	policy, decision-making rule
$\pi(s)$	action taken in state s under <i>deterministic</i> policy π
$\pi(a s)$	probability of taking action a in state s under <i>stochastic</i> policy π
$p(s', r s, a)$	probability of transitioning to state s' , with reward r , from s, a
$v_\pi(s)$	value of state s under policy π (expected return)
$v_*(s)$	value of state s under the optimal policy
$q_\pi(s, a)$	value of taking action a in state s under policy π
$q_*(s, a)$	value of taking action a in state s under the optimal policy
$V_t(s)$	estimate (a random variable) of $v_\pi(s)$ or $v_*(s)$
$Q_t(s, a)$	estimate (a random variable) of $q_\pi(s, a)$ or $q_*(s, a)$
$\hat{v}(s, \mathbf{w})$	approximate value of state s given a vector of weights \mathbf{w}
$\hat{q}(s, a, \mathbf{w})$	approximate value of state–action pair s, a given weights \mathbf{w}
\mathbf{w}, \mathbf{w}_t	vector of (possibly learned) <i>weights</i> underlying an approximate value function
$\mathbf{x}(s)$	vector of features visible when in state s

Appendix

Notation 3/3

δ_t	temporal-difference error at t (a random variable, even though not upper case)
$E_t(s)$	eligibility trace for state s at t
$E_t(s, a)$	eligibility trace for a state–action pair
\mathbf{e}_t	eligibility trace vector at t
γ	discount-rate parameter
ε	probability of random action in ε -greedy policy
α, β	step-size parameters
λ	decay-rate parameter for eligibility traces

David Silver's Slides

<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>

Lilian Weng's Blog

<https://lilianweng.github.io/lil-log/2018/02/19/a-long-peek-into-reinforcement-learning.html>

Other notations used in some fields

State	s	x
Action	a	u
Reward	r	$-c$